

DATA CLEANING

&

PREPROCESSING

in Machine Learning

**Submitted By:
Janeeta Ishtiaq**

Data Cleaning & Preprocessing in Machine Learning

Before training any machine learning model, data must be cleaned and transformed into a usable format. Raw data often contains **missing values, inconsistent formats, duplicates, outliers, and irrelevant information**. This article walks through important preprocessing techniques with Python examples.

1 Handling Missing Values

Missing values are very common in datasets (e.g., missing ages, salaries, or survey answers). If left untreated, they can bias results or break algorithms.

Common solutions include:

- **Dropping** missing rows or columns (only if data loss is small).
- **Filling** missing values with statistical measures like:
 - Mean (for continuous numerical data)
 - Median (when data is skewed)
 - Mode (for categorical data)

```
#1-Handling missing values
data1={
    "Name":["Janeeta","Ali","Sana","Hassan","Ahmed"],
    "Age":[22,np.nan,30,25,28],
    "Salary":[50000,40000,np.nan,80000,90000]
}
df1=pd.DataFrame(data1)
print(df1)

#Drop missing rows/columns:
# df.dropna(inplace=True)
# print(f"Drop missing rows/columns:\n {df}")

#Fill with mean/median/mode:
df1.fillna({"Age":int(df1["Age"].mean()),"Salary":int(df1["Salary"].median())},inplace=True)
print(f"Fill with mean/median/mode:\n {df1}")
```

Output:

```
      Name  Age  Salary
0  Janeeta  22.0  50000.0
1     Ali   NaN  40000.0
2     Sana  30.0     NaN
3  Hassan  25.0  80000.0
4  Ahmed  28.0  90000.0
Drop missing rows/columns:
      Name  Age  Salary
0  Janeeta  22.0  50000.0
3  Hassan  25.0  80000.0
4  Ahmed  28.0  90000.0
```

```
Fill with mean/median/mode:
      Name  Age  Salary
0  Janeeta  22.0  50000.0
1     Ali   26.0  40000.0
2     Sana  30.0  65000.0
3  Hassan  25.0  80000.0
4  Ahmed  28.0  90000.0
```

2. Data Type Conversion (Dates, Prices)

Datasets often have values stored in the wrong format:

- Dates stored as text ("2025-02-01") instead of **datetime**.
- Prices stored with symbols ("Rs500") instead of numbers.

Converting them to the right format ensures correct calculations, sorting, and analysis.
For example:

- Converting "Rs500" → 500
- Extracting day, month, year from "2025-02-01"

```
27 data2={
28     "Date":["2025-02-01","2025-03-21","2025-05-30","2025-07-01"],
29     "Price":["Rs500","Rs800","Rs250","Rs1200"]
30 }
31 df2=pd.DataFrame(data2)
32 print(df2)
33 df2["Date"]=pd.to_datetime(df2["Date"])
34 df2["Price"] = df2["Price"].replace('Rs', '', regex=True).astype(int)
35 print(f"After Conversion:\n{df2}")
36
```

Output:

```
      Date  Price
0  2025-02-01  Rs500
1  2025-03-21  Rs800
2  2025-05-30  Rs250
3  2025-07-01  Rs1200
After Conversion:
      Date  Price
0  2025-02-01    500
1  2025-03-21    800
2  2025-05-30    250
3  2025-07-01   1200
```

3 Text Cleaning (HTML, Stopwords)

When working with textual data (reviews, comments, web pages), the raw text usually contains:

- **HTML tags** (<p>,)
- **Special characters / punctuation** (!!!, ...)
- **Stopwords** (common words like *is*, *the*, *an*) that don't carry useful meaning.

Cleaning involves:

1. Removing HTML tags
2. Removing special characters
3. Converting to lowercase
4. Removing stopwords

This process makes text **simpler and meaningful** for NLP tasks like sentiment analysis.

```
37 #Text Cleaning(HTML,Stopwords)
38 text = "<p>Hello!!! This is <b>an example</b> sentence...</p>"
39 cleaned=BeautifulSoup(text,"html.parser").get_text()
40 cleaned=re.sub(r"[^a-zA-Z]", " ",cleaned)
41 cleaned=cleaned.lower()
42 stop_words = set(stopwords.words("english"))
43 cleaned = " ".join([word for word in cleaned.split() if word not in stop_words])
44 print(cleaned)
```

Output:



4 Normalization & Standardization

Machine learning models perform better when all features are on the **same scale**.

- **Normalization**: Rescales values between 0 and 1.
Useful when features have very different ranges (e.g., income vs age).
- **Standardization**: Rescales values so they have **mean = 0** and **standard deviation = 1**.
Useful for algorithms sensitive to distribution, like **SVM** or **K-Means**.

```
46 #Normalization & Standardization
47 data3=np.array([[150],[160],[170],[180],[190]])
48 df3=pd.DataFrame(data3,columns=["Heights"])
49 print(f"Original data:\n {df3}")
50 minmax_scaler = MinMaxScaler()
51 df3["Normalized"]=minmax_scaler.fit_transform(df3[["Heights"]])
52 standard_scaler = StandardScaler()
53 df3["standardized"]=standard_scaler.fit_transform(df3[["Heights"]])
54 print(f"After Normalization and Standaradization \n: {df3}")
55
```

Output:

```
##110 Example sentence
Original data:
      Heights
0        150
1        160
2        170
3        180
4        190
After Normalization and Standaradization
:      Heights  Normalized  standardized
0        150         0.00      -1.414214
1        160         0.25      -0.707107
2        170         0.50       0.000000
3        180         0.75       0.707107
4        190         1.00       1.414214
```

5 Duplicates & Outliers Handling

- **Duplicates:** Sometimes the same rows appear multiple times in data. Removing duplicates avoids bias and redundancy.
- **Outliers:** Extremely high or low values (e.g., salary = 1,000,000 in a dataset of 40k–90k).
 - Can distort averages and affect models.
 - Common approaches:
 - Remove outliers
 - Replace them with mean/median values

```
55
56 #Duplicates & Outliers Handling
57 data4={
58     "Name":["Ali","Sara","Ali"],
59     "Age":[22,30,22]
60 }
61 df4=pd.DataFrame(data4)
62 print(f"data with dupliacte rows:\n{df4}")
63 df4=df4.drop_duplicates(df4)
64 print (f"After handling duplicates :\n{df4}")
65
66 #Outliers
67 salary = [40000, 45000, 42000, 1000000]
68 q1, q3 = np.percentile(salary, [25, 75])
69 iqr = q3 - q1
70 lower = q1 - 1.5 * iqr
71 upper = q3 + 1.5 * iqr
72 outliers = [x for x in salary if x < lower or x > upper]
73 print(outliers)
74 salary_no_outliers = [x for x in salary if lower <= x <= upper]
75 mean_value = int(np.mean(salary_no_outliers))
76 salary_mean_replaced = [mean_value if (x < lower or x > upper) else x for x in salary]
77 print("Replaced with Mean:", salary_mean_replaced)
```

Output:

```
data with dupliacte rows:
```

```
   Name  Age
0  Ali   22
1  Sara  30
2  Ali   22
```

```
After handling duplicates :
```

```
   Name  Age
0  Ali   22
1  Sara  30
```

```
[1000000]
```

```
Replaced with Mean: [40000, 45000, 42000, 42333]
```

6 Encoding Categorical Data

Machine learning models work with **numbers**, not text.

Categorical variables (like *city = Lahore, Karachi*) must be converted:

- **Label Encoding:** Assigns each category a numeric value (e.g., *Lahore=0, Karachi=1*).
- **One-Hot Encoding:** Creates separate columns for each category (*City_Lahore, City_Karachi*).
 - Prevents models from assuming order in categories.

```
79 #Encoding Categorical Data
80 data5 = {"City": ["Lahore", "Karachi", "Lahore"]}
81 df5 = pd.DataFrame(data5)
82 #Label Encoding
83 le = LabelEncoder()
84 df5["City_encoded"] = le.fit_transform(df5["City"])
85 print(df5)
86 #One-Hot Encoding
87 df5 = pd.get_dummies(df5, columns=["City"])
88 print(df5)
```

Output:

	City	City_encoded
0	Lahore	1
1	Karachi	0
2	Lahore	1

	City_encoded	City_Karachi	City_Lahore
0	1	False	True
1	0	True	False
2	1	False	True

7 Feature Engineering

Feature engineering means **creating new features** from existing data to improve model learning.

Examples:

- **Word count:** Counting words in a review → helps in sentiment analysis.
- **Date parts:** Extracting year, month, weekday from a date → helps in time-series forecasting.

This step makes models smarter by adding **domain knowledge** into features.

```

91 #Feature Engineering
92 #Word Count (Text Data)
93 df = pd.DataFrame({
94     "Review": ["This product is good", "Worst experience"]
95 })
96
97 df["word_count"] = df["Review"].apply(lambda x: len(x.split()))
98 print(df)
99
100 #Date Parts (Date/Time Data)
101 df = pd.DataFrame({
102     "Order_Date": pd.to_datetime(["2024-01-15", "2024-06-30"])
103 })
104 df["year"] = df["Order_Date"].dt.year
105 df["month"] = df["Order_Date"].dt.month
106 df["day_of_week"] = df["Order_Date"].dt.dayofweek
107 print(df)

```

Output:

	Review	word_count
0	This product is good	4
1	Worst experience	2

	Order_Date	year	month	day_of_week
0	2024-01-15	2024	1	0
1	2024-06-30	2024	6	6

8 Merging Datasets

In real-world projects, data often comes from **multiple sources** (customers table, orders table, transactions table).

Merging datasets combines them into one complete dataset for analysis.

For example: joining **customer info** with **order history** using a common key (CustomerID).

```
109 #Merge Datasets
110 customers = pd.DataFrame({
111     "CustomerID": [1, 2],
112     "Name": ["Ali", "Sana"]
113 })
114
115 orders = pd.DataFrame({
116     "OrderID": [101, 102],
117     "CustomerID": [1, 2],
118     "Amount": [500, 300]
119 })
120
121 merged = pd.merge(orders, customers, on="CustomerID")
122 print(merged)
```

Output:

```
   OrderID  CustomerID  Amount  Name
0      101           1     500   Ali
1      102           2     300  Sana
```

Conclusion

Data cleaning & preprocessing are the **foundation of machine learning**.

We covered:

- ✓ Handling missing values
- ✓ Data type conversion
- ✓ Text cleaning
- ✓ Normalization & standardization

- ✓ Removing duplicates & outliers
- ✓ Encoding categorical data
- ✓ Feature engineering
- ✓ Merging datasets

The cleaner the data → the better the features → the more accurate the model

